

Chapter 12

Eligibility Traces





Eligibility Traces (적격성 추적)

- Eligibility Trace 와 TD를 함께 사용하면 학습을 더 효과적으로 수행할 수 있는 일반식을 만들 수 있음.
Q-Learning, Sarsa
General Method
- Eligibility Trace 로 TD, MC 를 통합하여 표현할 수 있음.
 - $\lambda=1$: Monte Carlo (method at one end)
 - $\lambda=0$: one-step TD



Eligibility Traces 의 특징

Eligibility Traces 의 특징

- Short-term memory vector, the eligibility traces $\mathbf{z}_t \in \mathbb{R}^d$
- Long-term weight vector $\mathbf{w}_t \in \mathbb{R}^d$

Advantages

- Only a single Trace Vector is required (자원)
 - Vs n feature vectors
- Learning occurs continually and uniformly in time (시간)
 - Vs delayed or end of the episode
- Learning can occur and affect behavior immediately (효율)
 - vs being delayed n steps





Forward views and Backward views

- ◎ Forward views
 - 앞의 보상을 모르고 update 하는 공식들.
 - 미래 보상(현재 유효하지 않음) 보기 때문에 구현 어려움.
- ◎ Backward views
 - 현재 TD Error 와 과거(backward) 방문 state 의 eligibility trace 사용



Forward views and Backward views







MC, TD and n-step TD

- MC : $G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots + \gamma^{T-t-1} R_T$
- TD : $G_t^{(1)} \doteq R_{t+1} + \gamma V_t(S_{t+1})$
- 2 step TD : $G_t^{(2)} \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 V_t(S_{t+2})$
- N step TD: $G_t^{(n)} \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n V_t(S_{t+n}) \quad (7.1)$
- N step TD with state and weight vector (12.1)
$$G_{t:t+n} \doteq R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n \hat{v}(S_{t+n}, \mathbf{w}_{t+n-1}), \quad 0 \leq t \leq T-n,$$



Weighting and decay


- TD(λ) 의 weighting :
 - 처음 것은 $1 - \lambda$
 - 이후 weight 합이 1이 되도록 λ 만큼 계속 곱하게 됨.
- Decay
 - 갈 수록 weight 는 0 수렴
- If..
 - $\lambda = 1$: MC
 - $\lambda = 0$: one-step TD





RMS comparison

- Off-line λ -return algorithm is slightly better at best α and λ , and at high α
- TD(λ) is worse at high α value





SARSA(λ)

From ch. 10,

$$\begin{aligned} G_{t:t+n} &\doteq R_{t+1} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n \hat{q}(S_{t+n}, A_{t+n}, \mathbf{w}_{t+n-1}), \quad t+n < T, \\ \mathbf{w}_{t+1} &\doteq \mathbf{w}_t + \alpha \left[G_t^\lambda - \hat{q}(S_t, A_t, \mathbf{w}_t) \right] \nabla \hat{q}(S_t, A_t, \mathbf{w}_t), \quad t = 0, \dots, T-1, \end{aligned}$$

Action value, TD error and eligibility trace for Sarsa(λ)

$$\mathbf{w}_{t+1} \doteq \mathbf{w}_t + \alpha \delta_t \mathbf{z}_t,$$

$$\delta_t \doteq R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}, \mathbf{w}_t) - \hat{q}(S_t, A_t, \mathbf{w}_t),$$


$$\mathbf{z}_{-1} \doteq \mathbf{0},$$

$$\mathbf{z}_t \doteq \gamma \lambda \mathbf{z}_{t-1} + \nabla \hat{q}(S_t, A_t, \mathbf{w}_t), \quad 0 \leq t \leq T.$$




SARSA(λ)


Path taken




Action values increased by one-step Sarsa



Action values increased by 10-step Sarsa



Action values increased by Sarsa(λ) with $\lambda=0.9$

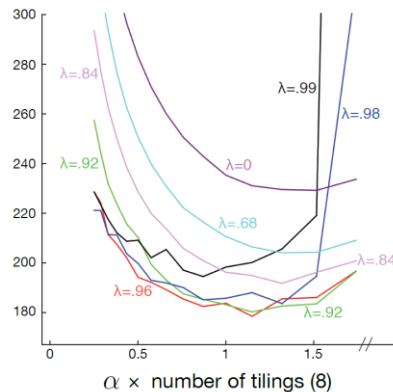


2. Sarsa(λ) @ ex 10.1


At Sarsa(λ),
more efficient learning.

Mountain Car
Steps per episode
averaged over
first 50 episodes
and 100 runs

Sarsa(λ) with replacing traces




n-step Sarsa

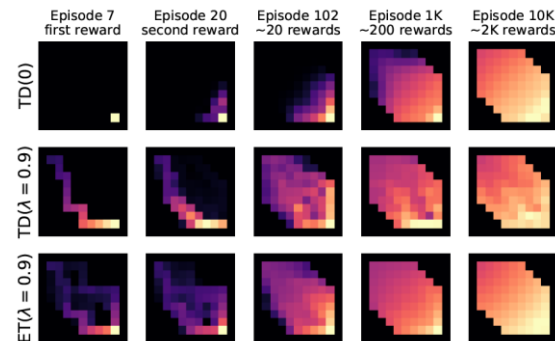




Performance



- TD(0) only updated the last state
 - TD(λ) updated the trajectory in this episode
 - ET(λ) additionally updated trajectories from earlier (unrewarding) episodes.
-
- TD(0) learns slowly but steadily,
 - TD(λ) learns faster but with higher variance
 - ET(λ) learns both fast and stable.





Performance




Figure 4: Prediction errors in the multi-chain. $ET(\lambda)$ (orange) consistently outperformed $TD(\lambda)$ (blue). Shaded areas depict standard errors across 10 seeds.




Figure 5: Comparing value error with linear function approximation a) as function of the number of branches (left), b) as function of λ (center two plots) and c) as function of η (right). The left three plots show comparisons of $TD(\lambda)$ (blue) and $ET(\lambda)$ (orange), showing $ET(\lambda)$ attained lower prediction errors. The right plot interpolates between these algorithms via $ET(\lambda, \eta)$, from $ET(\lambda) = ET(\lambda, 0)$ to $ET(\lambda, 1) = TD(\lambda)$, with $\lambda = 0.9$ (corresponding to a vertical slice indicated in the second plot).



Thanks!

Any **questions** ?